

Console Open AI API chat

<Implement your console chat with Open AI API>






Request

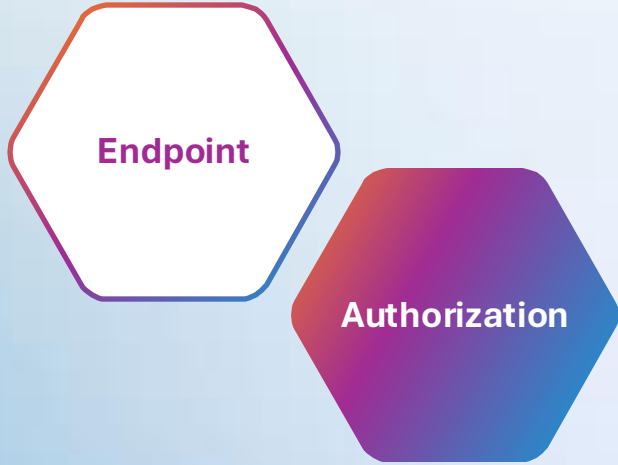
<Request and its params>

Request:



<https://api.openai.com/v1/chat/completions>

Request:



Authorization: Bearer {YOUR_API_KEY}

<https://platform.openai.com/settings/organization/api-keys>

Request:

Endpoint

Authorization

Request body

```
{  
  "model": "gpt-4o",  
  "messages": [  
    {  
      "role": "system",  
      "content": "You are a helpful assistant."  
    },  
    {  
      "role": "user",  
      "content": "Write a haiku about recursion in  
programming."  
    }  
  ],  
  "stream": false  
}
```

Request:

Endpoint

Authorization

Request body

```
{  
  "model": "gpt-4o",  
  "messages": [  
    {  
      "role": "system",  
      "content": "You are a haiku writer."  
    },  
    {  
      "role": "user",  
      "content": "Write a haiku about recursion in  
programming."  
    }  
  ],  
  "stream": false  
}
```

Request:


Endpoint

Authorization

Request body

```
{  
  "model": "gpt-4o",  
  "messages": [  
    {  
      "role": "system",  
      "content": "You are a helpful assistant."  
    },  
    {  
      "role": "user",  
      "content": "Write a haiku about recursion in  
programming."  
    }  
  ],  
  "stream": false  
}
```




Request:



Enables response streaming, where the server sends back partial responses as they are generated. This is useful for applications needing real-time updates.




Default value: *false*

Request:






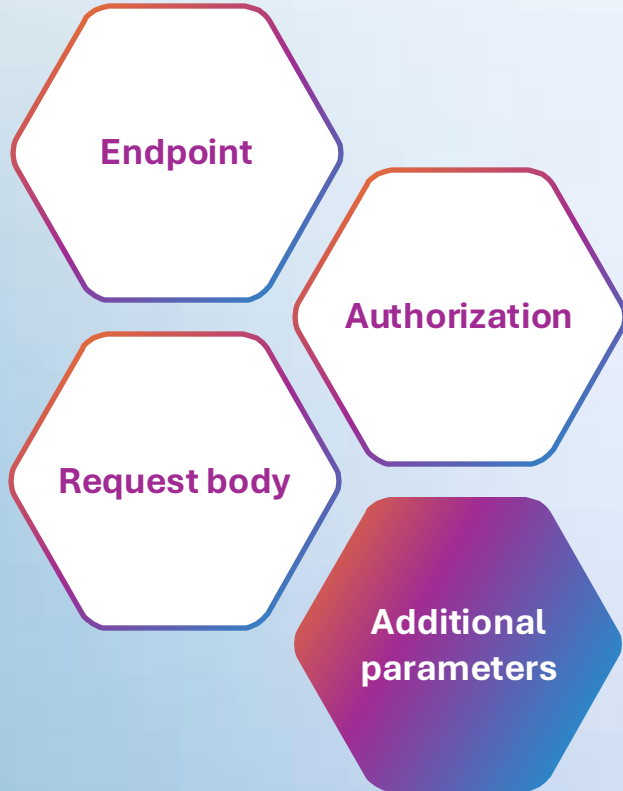
Controls the randomness of the output.
It's a parameter for balancing creativity and determinism.
Range: *0.0* to *2.0*
Default value: *1.0*

Request:



An alternative to “*temperature*” for controlling diversity
Via nucleus sample.
Range: *0.0* to *1.0*
Default value: *1.0*
USE “*top_p*” OR “*temperature*” NOT BOTH


Request:



Sets the maximum number of tokens to generate in the response.

Default value: ---

Request:



stream

temperature

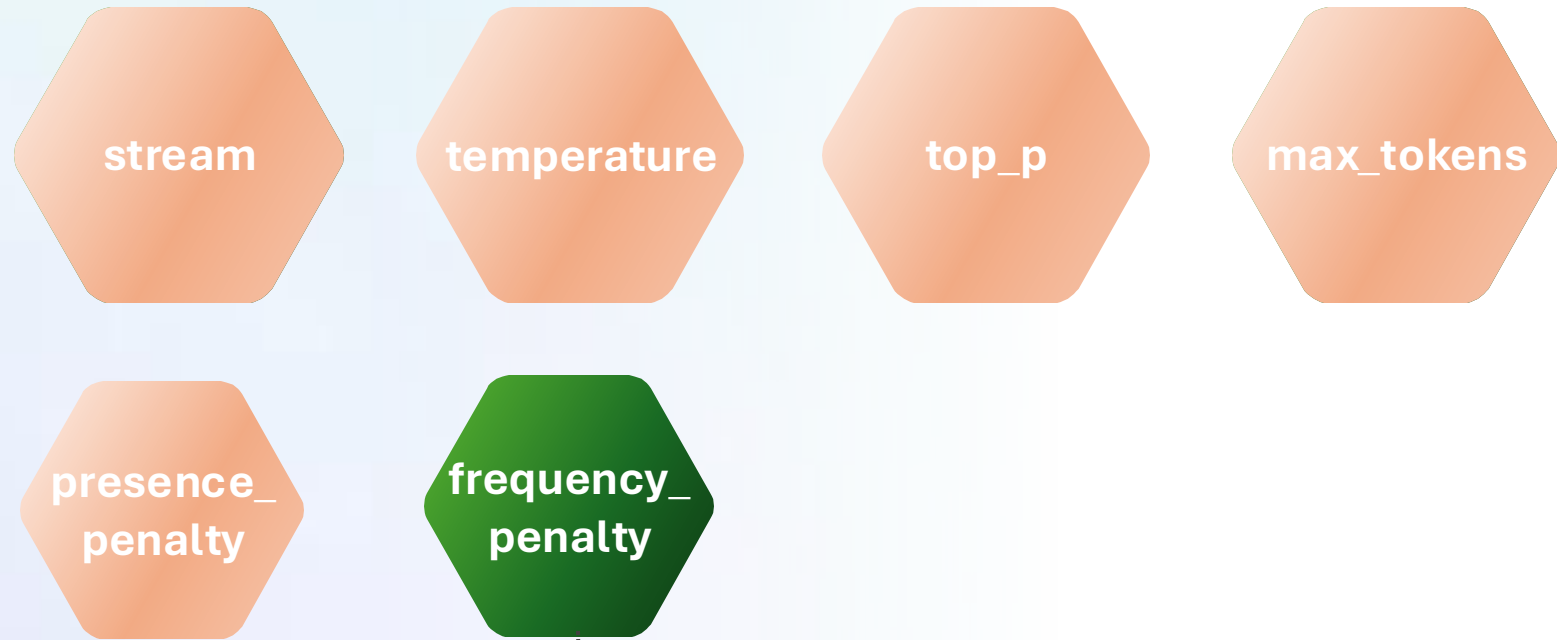

top_p

max_tokens

presence_penalty

Penalizes new tokens based on whether they appear in the text so far, encouraging the model to discuss new topics.
Range: **-2.0** to **2.0**
Default value: **0.0**
Positive: *Drives novelty and creativity.*
Negative: *Keeps the model focused on familiar ideas, potentially leading to more repetitive outputs.*

Request:



Penalizes new tokens based on their frequency in the text so far, reducing verbosity or repetitiveness.

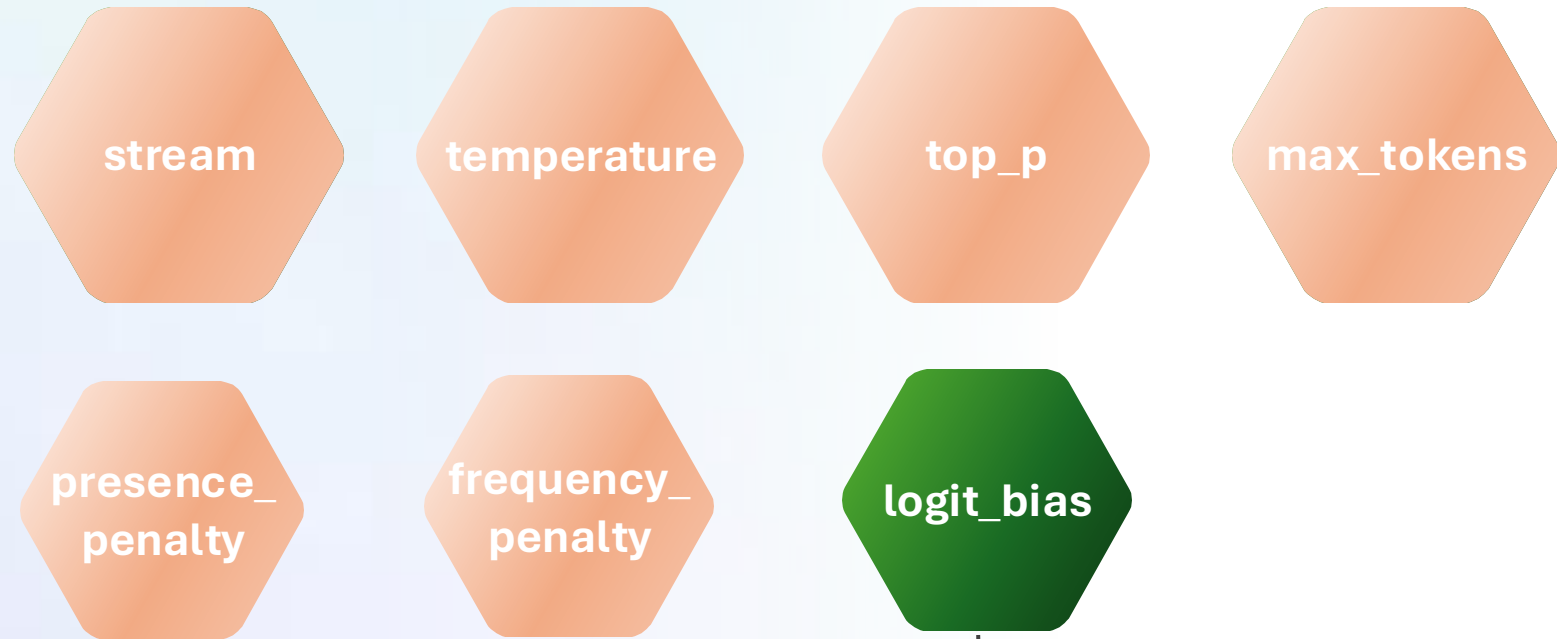

Range: -2.0 to 2.0

Default value: 0.0

Negative values encourage more repetition.

Positive values reduce token repetition.

Request:



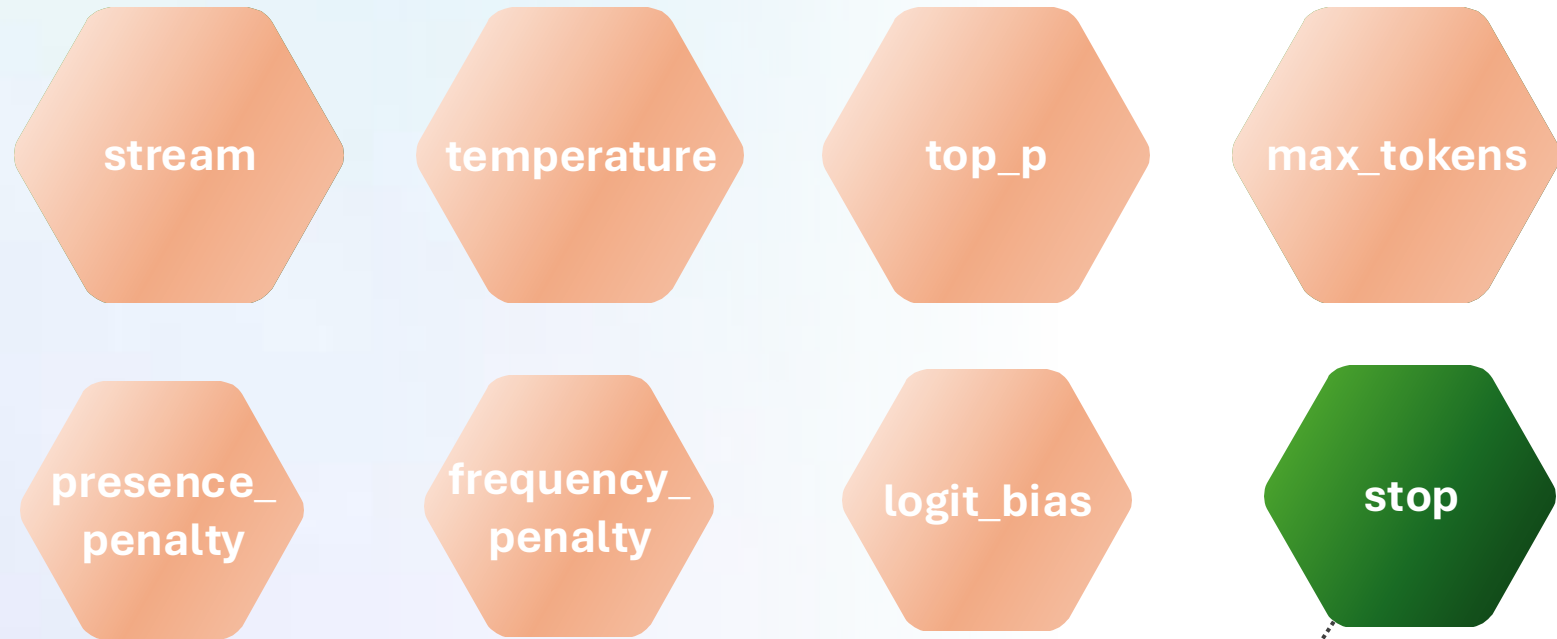

Adjusts the probability of specific tokens appearing in the response.

`Map>string, number< ("and": +100)`
Default value: *null*

A map of token IDs to bias values (e.g., -100 to prevent a token, +100 to force a token).

Useful for controlling specific behaviors.

Request:




Specifies one or more sequences where the API should stop generating further tokens.

**[“\n”, “codeus”]
Default value: *null***

***Specifies one or more sequences where the API should stop generating further tokens.
Can include multiple stop conditions.***

Request:



stream

temperature

top_p

max_tokens

presence_penalty

frequency_penalty

logit_bias

stop

n

Specifies amount of assistant messages that should be generated.

Default value: *1*



Response

<Response and its handling>

Response:

Default

```
{
  "id": "chatcmpl-7vHVghsEnUygFKP9sBxX3TEBh4L3q",
  "object": "chat.completion",
  "created": 1678928387,
  "model": "gpt-4",
  "usage": {
    "prompt_tokens": 12,
    "completion_tokens": 24,
    "total_tokens": 36
  },
  "choices": [
    {
      "index": 0,
      "message": {
        "role": "assistant",
        "content": "Why don't scientists trust atoms? Because they make up everything!"
      },
      "finish_reason": "stop"
    }
  ]
}
```

Response:

Default

```
{
  "id": "chatcmpl-7vHVghsEnUygFKP9sBxX3TEBh4L3q",
  "object": "chat.completion",
  "created": 1678928387,
  "model": "gpt-4",
  "usage": {
    "prompt_tokens": 12,
    "completion_tokens": 24,
    "total_tokens": 36
  },
  "choices": [
    {
      "index": 0,
      "message": {
        "role": "assistant",
        "content": "Why don't scientists trust atoms? Because they make up everything!"
      },
      "finish_reason": "stop"
    }
  ]
}
```

Response:

Default

```
{
  "id": "chatcmpl-7vHVghsEnUygFKP9sBxX3TEBh4L3q",
  "object": "chat.completion",
  "created": 1678928387,
  "model": "gpt-4",
  "usage": {
    "prompt_tokens": 12,
    "completion_tokens": 24,
    "total_tokens": 36
  },
  "choices": [
    {
      "index": 0,
      "message": {
        "role": "assistant",
        "content": "Why don't scientists trust atoms? Because they make up everything!"
      },
      "finish_reason": "stop"
    }
  ]
}
```

Response:

Default

Streaming
stream: true

Chunk 1

```
data: {  
  "choices": [  
    {  
      "delta": {  
        "role": "assistant",  
        "content": "Why don't scientists"      }  
    }  
  ]  
}
```

Response:

Default

Streaming
stream: true

Chunk 1

```
data: {  
  "choices": [  
    {  
      "delta": {  
        "role": "assistant",  
        "content": "Why don't scientists"      }  
    }  
  ]  
}
```

Response:

Default

Streaming
stream: true

Chunk 1

Chunk 2

```
data: {  
  "choices": [  
    {  
      "delta": {  
        "content": " trust atoms?"  
      }  
    }  
  ]  
}
```

Response:

Default

Streaming
stream: true

Chunk 1

Chunk 2

Chunk 3

```
data: {  
  "choices": [  
    {  
      "delta": {  
        "content": " Because they make up  
everything!"  
      }  
    }  
  ]  
}
```


Response:

Default

Streaming
stream: true

Final chunk

Chunk 1

Chunk 2

Chunk 3

data: [DONE]

Thank you

- Author: Pavlo Khshanovskyi
- My LinkedIn: Link to your LI
- Date: January 2025
- [Join Codeus community in Discord](#)
- [Join Codeus community in LinkedIn](#)